



Reactor

Model Mondays

<https://aka.ms/model-mondays>

JOIN THE COMMUNITY

<https://aka.ms/model-mondays/discord>





Reactor

MODEL MONDAYS S4:E01

Spotlight on MAI Models



Hosted By: Amy Boyd

Aug 10, 2026 · 1:30-2:30 PM ET

Sophie
Lebrecht

Member of
Technical Staff
Microsoft AI



Yanan
Cai

Product Mgr.
Microsoft AI





Reactor

PART 1: OVERVIEW

Microsoft AI (MAI) Models

Sophie Lebrecht

Member Of Technical Staff, Microsoft AI



The curve is shifting
From *tokenmaxxing* to *token efficiency*

- 1 AI Infra
- 2 Silicon
- 3 Reasoning
- 4 Coding
- 5 Tool use
- 6 Helpfulness
- 7 Safety
- 8 Efficiency
- 9 Harnesses
- 10 RLEEs



MAI models are built from scratch, not distilled



Trained on a clean lineage of commercially safe training data



Designed to ensure more efficient Return on Token spend



Continuously improving through our hill climbing machine

Powering billions. Now ready to power what you build.



Excel

Equally performant
but more cost-efficient
than GPT-5.6 for most common
tasks



GitHub

10%
higher code accept rate and
lower median token usage
than GPT-5.4 Mini and Claude
Haiku 4.5 in VS Code



PowerPoint

84%
reduced GPU costs
compared to GPT-Image-2



Dragon Copilot

50%
relative reduction
in transcription and language-
identification error rates



Dynamics 365

89%
reduced GPU costs
for customers like
T-Mobile and EasyJet

Providing the right model for the right job

An efficient multimodal stack across chat, image, voice, and speech



Reasoning

MAI-Thinking-1
MAI-Code-1-Flash



Image

MAI-Image-2.5 Pro
MAI-Image-2.5
MAI-Image-2.5-Flash



Audio

MAI-Transcribe-1.5
MAI-Voice-2
MAI-Voice-2-Flash



Cyber

MAI-Cyber-1-Flash

Hill-climbing with your models




Figure 1:

MAI-Code-1-Flash, post-trained inside the GitHub Copilot harness, became the starting checkpoint for climbing on Excel evaluations—producing two highly efficient, specialized models from one lineage.

The Excel climb






Figure 2:

The MAI-Code-1-Flash checkpoint trained inside an Excel reinforcement learning environment, learning spreadsheet tools and knowledge workflows.

The advantage isn't the model. It's owning the loop that improves it.



Get started today



Microsoft Foundry

Primary enterprise API surface



aka.ms/MAlonFoundry



MAI Playground

Low-friction access to test models



aka.ms/MAlplayground



OpenRouter

Third-party hubs

Available through OpenRouter



aka.ms/MAlonOpenRouter



Reactor

PART 2: DEEP DIVE

MAI Coding Models

Yanan Cai

Product Manager, Microsoft AI



Built for GitHub Copilot

Trained and measured inside the real GitHub developer surface

Training

- GitHub Product Harness
 - VS Code
 - Copilot CLI

Evaluation

- Public and internal benchmarks
- In production metrics on **Quality, Efficiency** and **Engagement**

Controlled A/B experiment

Three signal families we measure on live GitHub Copilot traffic

Code quality

- How much code gets accepted
- How much code is retained after x mins
- How much code survives through a git commit

Token efficiency


- How many tokens needed to get a task done
- How many turns users need to reach a commit

Engagement

- Are users more likely to return and use it
- How many user-initiated turns

Hill-climbing with product experiments

Every gain proven on live product signals, not public benchmarks.



← → Demo1

EXPLORER

...

DEMO1

index.html

script.js

styles.css

Open Chat

Show All Commands

Start Debugging

^ ⌘ I

⇧ ⌘ P

F5

CHAT

FICTIONAL REDMOND COFFEE ROASTER WEBSITE

A website for a fictional Redmond coffee roaster

I'm setting up a small website for the fictional Redmond coffee roaster and will create the project structure and landing page in the workspace.

Reviewed Demo1 and created index.html, styles.css, script.js

Read

Demo1

Created index.html

Created styles.css

Created script.js

The landing page is in place; I'm checking that it renders cleanly and that the files are ready to open in a browser.

Ran `python3 -m http.server 8000 >/tmp/redmond_roast_se...` in sandbox

Ran `python3 -m http.server 8000 >/tmp/redmond_roast_se...`

Opened `Coding Benchmarks`

Running `python3 -m http.server 8080` in background — [Show](#)

Opened `Redmond Roast Co.`

> 3 files changed +1021 -3

Keep

Undo

Describe what to build

+

Agent

MAI-Code-1.1-Flash

Medium

Local [Allow all](#)

0 0

memos

EXPLORER

MEMOS

.github

cmd

docs

internal

proto

scripts

server

store

web

.dockerignore

.gitignore

.golangci.yaml

.release-please-mani...

AGENTS.md

CHANGELOG.md

CODEOWNERS

go.mod

go.sum

LICENSE

memos_prod.db

README.md

release-please-confi...

SECURITY.md

OUTLINE

TIMELINE

main*

Memos

http://localhost:3001/

Memos

Home

August 2026

S	M	T	W	T	F	S
26	27	28	29	30	31	1
2	3	4	5	6	7	8
9	10	11	12	13	14	15
16	17	18	19	20	21	22
23	24	25	26	27	28	29
30	31	1	2	3	4	5

VIEWS

Tasks

test

Any thoughts...

+ Private

Save

32 minutes ago

todo today: MAI-Code-1.1-Flash

32 minutes ago

todo today: Model Monday

CHAT

ADD CATEGORY TO NOTE CARD MENU

in the app, for each note card, can you add a catagory to color context menu? have the name of the catagory editable and configurable. the colors and name maping should be gobally the same

The MCP server [linear-server](#) may have new tools and requires interaction to start. [Start it now?](#)

I'm tracing the note card color context menu and the existing category/color model so I can implement the shared editable mapping cleanly and verify it in the app.

Searched for multiple regex patterns across codebase

I've narrowed this to the memo UI and the shared tag metadata settings; I'm checking whether the app already has a memo color model I can extend without introducing a second, inconsistent mapping.

Searched for various regex patterns and reviewed multiple files

I've confirmed there isn't an existing memo category model in the app, so I'm adding a small shared category-color registry with tests first, then I'll wire it

Add memo category UI and storage (2/4)

Memos

Describe what to build

Agent MAI-Code-1.1-... Medium

Local Allow all

usememos/memos: Open-sourcetonybaloney/vscode-pets: Ad

New Tab

https://github.com/tonybaloney/vscode-pets

UI ServiceBuild infoProduct infoDev toolsAdminPerf

tonybaloney / vscode-pets

Type / to search

CodeIssues36Pull requests40AgentsActionsSecurity and qualityInsights

Public

SponsorWatch9Fork538Star4.2k

main25 Branches17 TagsGo to fileCode

dependabot[bot] Bump actions/setup-node from 5 to 7 ...2c91214 · 3 days ago1,519 Commits

.github	Bump actions/setup-node from 5 to 7 (#928)	3 days ago
.vscode	chore: Update deprecated tsc.autoDetect setting...	2 months ago
docs	"Spawn Pet" instead of "Spawn Additional Pet" (...)	10 months ago
l10n	Add GUI button for changing theme (#891)	2 weeks ago
media	Add GUI button for changing theme (#891)	2 weeks ago
src	Add GUI button for changing theme (#891)	2 weeks ago
.eslintrc.json	Clean up all the dangling promises in the extension	3 years ago
.gitignore	Gitignore .DS_Store	3 years ago
.prettierrc	Incorporate no unused variables and no unused i...	4 years ago
.prettierrc.js	Implemented a Export Pet Command that export...	4 years ago
.vscodeignore	Change version and ignore node_modules	4 years ago

About

Adds playful pets 🐶🐱🐹 in your VS Code window

tonybaloney.github.io/vscode-pets

hacktoberfesthacktoberfest-acceptedtypescriptvisual-studio-codevscodevscode-extension


ReadmeMIT licenseCode of conductContributingActivity4.2k stars9 watching538 forksReport repository

https://github.com/tonybaloney/vscode-pets/pulls


	MAI-Code-1.1-Flash		MAI-Code-1-Flash		Haiku 4.5		GPT 5.4 mini	
	Pass rate	Tokens usage	Pass rate	Tokens usage	Pass rate	Tokens usage	Pass rate	Tokens usage
SWE-Bench Verified	72.6	8.6K	71.6	10.8K	69.8	20.9K	69.2	9.4K
Terminal Bench 2.1	62.9	17.0K	51.7	14.2K	49.4	25.5K	60.7	21.9K

MAI-Code-1-Flash is
optimized for GitHub
Copilot in VS Code

Coding task reasoning
Agentic execution
Broad ecosystem support



Plan and reason through complex coding tasks from start to finish
Make decisions and adapt autonomously in multi-step workflows
Fluent across programming languages, frameworks, ecosystems



1 Try · MAI-Code-1-Flash



Reactor

PART 3: TRY OUT

MAI Image 2.5 Demo

Amy Boyd

Developer Relations - Microsoft



Providing the right model for the right job

An efficient multimodal stack across chat, image, voice, and speech



Reasoning

MAI-Thinking-1
MAI-Code-1-Flash



Image

MAI-Image-2.5 Pro
MAI-Image-2.5
MAI-Image-2.5-Flash



Audio

MAI-Transcribe-1.5
MAI-Voice-2
MAI-Voice-2-Flash




Cyber

MAI-Cyber-1-Flash

MAI-Image-2.5 Family

Create high-quality, photorealistic, and design ready images from simple text prompts or uploaded images



MAI

Models Blog Careers

THE HUMANIST REVIEW OF AI ↗

Model Family



MAI-Image-2.5-Pro

Robust object consistency with better visual reasoning and world knowledge.

Try in Playground →



MAI-Image-2.5

Built for high-quality generation and precise, controllable editing.

Try in Playground →



MAI-Image-2.5-Flash

Production-ready quality at less than half the price of our flagship model.

Try in Playground →

Read: <https://microsoft.ai/models/mai-image-2-5/>

Weekly Highlights

WHAT'S NEW IN
MICROSOFT FOUNDRY?


Sharmila Chockalingam

Director, AI Product Marketing @Microsoft



AT&T & Microsoft scale
trillion token workloads

Foundry Managed
Compute offers a
streamlined way to
access dedicated GPU



Using Microsoft Foundry Managed Compute, AT&T was able to experiment across multiple open models, optimize workloads across different GPU architectures, and process massive volumes of telecom data all within a unified platform

July 23 • 5 min read

AT&T and Microsoft scale trillion-token workloads with Microsoft Foundry and AMD

By [Steve Sweetman](#), Vice President of Product Management, Foundry Models

Listen to this post



Powered by Microsoft Copilot

How Microsoft Foundry
helped power AT&T's


OTel2.0

Metric	Value
OTel 1.0 Downloads	Over 25M
GPUs Used Through Foundry Managed Compute	About 530
Tokens Processed for OTel2.0	About 1T
Tokens Trained for OTel2.0	About 400 B
Models used to train OTel	Phi-4, OSS 120B, Gemma 4

See: [AT&T and Microsoft scale trillion-token workloads with Microsoft Foundry and AMD](#) | Microsoft Azure Blog

GPT-transcribe and GPT-live-transcribe in Microsoft Foundry

New developer options for
transcription in stored audio
and live voice interactions



GPT-transcribe is the highest-accuracy ASR model from OpenAI · GPT-live-transcribe is designed for low-latency streaming transcription using the Realtime API with “tunable latency” option for developers

**gpt-transcribe**

[Details](#) [Benchmarks](#) [Responsible AI](#) [License](#)

A new real-time speech-to-text (STT) model with enhanced transcription accuracy and low-latency streaming capabilities.

 Azure OpenAI  Direct from Azure Version: 2026-07-28

[Catalog](#) / [Models](#) / **gpt-live-transcribe**

[←](#) **gpt-live-transcribe** Version: 2026-07-28

[Details](#) [Benchmarks](#) [Responsible AI](#) [License](#)

A new real-time speech-to-text (STT) model with enhanced transcription accuracy and low-latency streaming capabilities.


 Azure OpenAI  Direct from Azure Version: 2026-07-28



A New Generation of AI Voice Models from Open AI

Introducing Kimi K3 through Fireworks AI on Microsoft Foundry

Kimi K3 is an open-weight model from Moonshot AI with 2.8T parameters and 1M token context window



This is Moonshot AI's most capable model. Use it for long-horizon coding (sustained engineering tasks, large repositories) · Supports native vision for workflows that combine software engineering with visual reasoning ·

Model Card: <https://ai.azure.com/catalog/models/FW-Kimi-K3>

MICROSOFT FOUNDRY BLOG 2 MIN READ

Introducing Kimi K3 through Fireworks AI on Microsoft Foundry



davevoutila MICROSOFT

Jul 28, 2026

Need to process entire codebases, lengthy legal documents, or hours of meeting transcripts in a single prompt? Kimi K3's 1 million-token context window—one of the largest available—is now deployable on Microsoft Foundry through Fireworks AI. Here's what that massive context means for your applications.

What makes Kimi K3 Different

Kimi K3 is an open-weight model from Moonshot AI featuring 2.8 trillion parameters—making it one of the largest open models available. Its 1 million-token context window means you can process approximately 750,000 words (roughly 10 full novels or an entire codebase) in a single request. Fireworks AI provides the optimized inference infrastructure, while Microsoft Foundry handles enterprise deployment and governance.

Use Cases

- **Full repository code analysis:** Analyze entire codebases for security vulnerabilities or refactoring opportunities. For example, you could pass your entire Node.js application (package.json, all source files, and dependencies) in a single prompt and ask Kimi K3 to identify unused imports, potential SQL injection points, or suggest architectural improvements—no chunking or retrieval augmentation required.
- **Long-document summarization:** Process legal contracts, research papers, or compliance documents without chunking
- **Extended conversation memory:** Build chatbots that remember entire conversation histories


See: <https://techcommunity.microsoft.com/blog/azure-ai-foundry-blog/introducing-kimi-k3-through-fireworks-ai-on-microsoft-foundry/4540187>

3

Try · Kimi K3 (Fireworks)

Anthropic's most advanced Opus model for Enterprise Agents and Complex Work

Manage long-running agents that can handle complex workflows



Key model capabilities of Claude Opus 5

1M token native context window · Adaptive thinking (default on) · Support image, pdf and text input · Advanced tool use with search and programmatic tool calling · Use for coding, agent pipelines, enterprise workflows, financial analysis, browser automation


Model Card: <https://ai.azure.com/catalog/models/claude-opus-5>



**Claude Opus 5 is available
in Microsoft Foundry**

Foundry Labs is the home
for cutting-edge AI
experiments from Microsoft

Tour the refreshed Foundry
Labs experience today!



CODE & SOFTWARE ENGINEERING

Fara1.5

Next-Gen Computer-Use Model Family

A family of three small computer-use models — 4B, 9B, and 27B — built for browser automation. The 9B flagship sets SOTA among small computer-use models on Online-Mind2Web, nearly doubling Fara-7B.


[Try it now →](#)[Learn more](#)

GEOSPATIAL & EARTH SCIENCE

Aurora 1.5

Earth-system Foundation Model

Aurora 1.5 is Microsoft's open Earth-system foundation model that combines 22 new weather variables, hourly forecasting, and ensemble forecasting to deliver more accurate weather and climate intelligence.

[Try it now →](#)[Learn more](#)

CREATIVE & GENERATIVE MEDIA

Phi-4-Reasoning-Vision-15B

Compact Multimodal Reasoning Model

A compact open-weight multimodal reasoning model built on Phi-4-Reasoning + SigLIP-2 with a mid-fusion architecture. Hybrid reasoning design (THINK / NOTHINK modes), high-resolution perception with up to 3,600 visual tokens — competitive with models 10× its size.

[Try it now →](#)[Learn more](#)

BIOMEDICAL SCIENCES

RosettaFold3

Unified Biomolecular Structure Prediction

A unified biomolecular modeling system that predicts 3D structures of proteins, nucleic acids, and small molecules within a single framework. Combines multimodal transformers and generative diffusion; supports protein-ligand, protein-DNA, protein-RNA interactions with atom-level conditioning.

[Try it now →](#)[Learn more](#)

View: <https://labs.ai.azure.com/>

5 Explore · Foundry Labs



Reactor

MODEL MONDAYS S4:E02

Cohere for Agentic AI



Hosted By: Bethany Jepchumba

Aug 17, 2026 · 1:30-2:30 PM ET



Andy Tran

Principal Solution Architect / Cohere

WATCH THE
LIVESTREAM



Mon 1:30pm ET
Register Here

<https://aka.ms/model-mondays/RSVP>

JOIN THE
COMMUNITY



Fri 1:30pm ET
Join The Discord

<https://aka.ms/model-mondays/Discord>

STAR US ON
GITHUB



Explore Blogs,
Recaps & Schedules

<https://aka.ms/model-mondays>