



Reactor

Model Mondays

<https://aka.ms/model-mondays>

JOIN THE COMMUNITY

<https://aka.ms/model-mondays/discord>





Reactor

MODEL MONDAYS S4:E04

From Research To Reality



Hosted By: Nitya Narasimhan

Aug 31, 2026 · 1:30-2:30 PM ET

Gulsimo
Osimi

Technical PM,
Microsoft



Yash
Lara

Principal PM
Microsoft





Reactor

RESEARCH TO REALITY

Spotlight on Foundry Labs

Gulsimo Osimi

Technical Program Manager
Microsoft



Microsoft Foundry Labs | Early-Stage AI

+

https://labs.ai.azure.com

Chat

Foundry Labs

AboutInnovationsStoriesCommunity



Foundry Labs

The home for cutting-edge AI experiments from Microsoft

Innovations

Breakthrough AI you can try today

GENERAL PURPOSE

Microsoft Agent Framework

Animations

Open-Source Engine for Agentic AI Apps

tokensvision_encoderreasoning



CREATIVE & GENERATIVE MEDIA



Phi-4-Reasoning-Vision-15B

Compact Multimodal Reasoning Model



Microsoft Foundry Labs is the hub for AI innovation from Microsoft

The latest experiments ready for you to explore ... and build on





Reactor

RESEARCH TO REALITY

Spotlight on Fara 1.5

Yash Lara

Principal Program Manager
Microsoft




Fara1.5

A family of frontier computer use agent models for the browser — 4B, 9B and 27B, open weights under MIT.

Fara 1.5 Models

(4B, 9B, 27B variants)

Family of frontier
computer use agent
models for the browser
from Microsoft Research



Fara1.5-9B is a multimodal web agent developed by Microsoft Research AI Frontiers. It **observes the browser via screenshots and acts on behalf of the user** by emitting structured tool calls (*click*, *type*, *scroll*, *web_search*, *visit_url*, and *others*) to complete multi-step web tasks end-to-end.

A screenshot of the Microsoft Foundry web interface for the Fara1.5-9B model. The page has a dark theme with a grid pattern. At the top, there's a navigation bar with 'Enterprise', 'Discover' (selected), 'API reference', 'Pricing', and a 'Sign in' button. Below the navigation bar, the breadcrumb trail shows 'Catalog / Models / Fara1.5-9B'. The main heading is 'Fara1.5-9B' with a back arrow and the Microsoft logo. There are tabs for 'Details' (active), 'Benchmarks', 'Responsible AI', and 'License'. A description states: 'Multimodal web agent for browser-based task automation.' Below this, there's a 'Microsoft' logo and 'Version: 1'. The main content area has two paragraphs. The first paragraph describes the model's capabilities: 'Fara1.5-9B is a multimodal web agent developed by Microsoft Research AI Frontiers. It observes the browser via screenshots and acts on behalf of the user by emitting structured tool calls (click, type, scroll, web_search, visit_url, and others) to complete multi-step web tasks end-to-end. Primary use cases include filling forms, shopping, booking travel and restaurants, information seeking, and account-driven workflows. Fara1.5-9B is trained to recognize "critical points" in a task (situations involving missing user information, ambiguous instructions, or irreversible actions such as completing a purchase or sending a message) and to pause for user confirmation rather than proceed unilaterally. It is intended for human-in-the-loop, sandboxed deployment; the recommended deployment vehicle is MagenticLite, which provides allow-listed navigation, watch-mode action monitoring, an immediate pause control, and a Docker-isolated browser environment.' The second paragraph states: 'Fara1.5-9B is a 9B-parameter multimodal decoder-only language model fine-tuned from the Qwen 3.5 9B base. It accepts the user's textual goal, current browser screenshot(s), and a textual history of prior thoughts and actions, and emits a chain-of-thought block followed by a structured tool-call block. The model is vision-only, perceiving the browser exclusively through UI screenshots, and supports up to a 262K-token context window, which enables long multi-step trajectories with substantial screenshot and action-history accumulation.' On the right side, there's a 'Quick facts' section with a table: 'Model provider' (Microsoft), 'Type' (Web agent tasks, Gui grounding, Image to text), 'Lifecycle' (Preview), 'Input type' (image, text), and 'Output type' (text). Below this is a 'Quick start' section with three buttons: 'Sign in to get started', 'View documentation' (with an external link icon), and 'Azure support' (with an external link icon).

1 Try It · Fara 1.5 On Foundry

Visit: <https://ai.azure.com/catalog/models/Fara1.5-9B>



Reactor

DEVELOPER CORNER

Build with Foundry Canvas

Junjie Li



Senior Product Manager
Microsoft





Microsoft Foundry Canvas is a GitHub Copilot App extension

Design, configure, test, and deploy, Microsoft Foundry hosted agents without leaving your editor!



1 Try It · Foundry Canvas



MICROSOFT
FOUNDRY
DISCORD



Foundry Friday AMA

BUILD AGENTS WITH FOUNDRY
CANVAS & GITHUB COPILOT APP



Hosted By: Carlotta Castelluccio

Sep 4, 2026 · 10:00-10:30 AM CEST



Junjie Li

Senior Product Manager
Microsoft



Reactor

LAST WEEK IN FOUNDRY

News Highlights


Sharmila Chockalingam

Director, AI Product Marketing Microsoft



Cohere Parse v5

A Vision Language Model from Cohere is now available in Microsoft Foundry



Cohere Parse takes complex, multimodal files and returns clean Markdown suitable for downstream indexing and application logic.

MICROSOFT FOUNDRY BLOG 1 MIN READ

Introducing Cohere Parse V5



RashaudSavage MICROSOFT

Aug 27, 2026

Cohere Parse, a vision language model built for converting enterprise documents into structured, machine-readable data, is now available in the Microsoft Foundry. Parse joins Cohere's Embed and Rerank models in Foundry.

What Parse does

Cohere Parse takes complex, multimodal files and returns clean Markdown suitable for downstream indexing and application logic.

According to Cohere, Parse can do the following:

Goes beyond text recognition - Parse detects and interprets tables, forms, diagrams, and embedded images, extracting semantic context rather than a flat character stream.

Preserves document structure- The model returns bounding boxes for visual elements, which supports retrieval, grounding, and citation in downstream applications.

Is trained on enterprise document types- Cohere trained Parse on business documents across domains including finance, insurance, and scientific research and supports nine major world languages.


Is built for high-throughput workloads- Cohere reports a measured rate of 4.5 pages per second per GPU, or approximately 36 pages per second on an 8× H100 node, served with vLLM.

1

Try It · Cohere Parse v5

Grok 4.6

Long horizon reasoning
and complex workflows



Grok 4.6 focuses on strong performance across coding, engineering, office productivity, research enablement, and inference optimization tasks

MICROSOFT FOUNDRY BLOG 3 MIN READ

Grok 4.6 comes to Microsoft Foundry Models: Built for long-horizon reasoning and complex workflows



Mahesh Balachandran MICROSOFT


Aug 26, 2026

The next wave of AI applications won't be defined by how quickly models answer questions. They'll be defined by how much work they can complete. Whether resolving repository issues, navigating codebases, generating engineering designs, or executing multi-step enterprise workflows, developers are increasingly building systems that need models capable of staying on task across complex, long-running objectives.

Today, Grok 4.6 from SpaceXAI is available in Microsoft Foundry Models through public preview, bringing SpaceXAI's latest frontier model to developers through a unified platform for model discovery, evaluation, deployment, and governance.

Built on SpaceXAI's 1.5T-scale model family, Grok 4.6 focuses on s enablement, and inference optimization tasks. Rather than optimi execution across software engineering, agentic workflows, and tex

Benchmark highlights



2

Try It · Grok 4.6


LLM Evaluation

Evaluation led
development

Sharing best practices to move LLMs from prototype to production in code analysis, developer tools, security, data analysis and other production workflows

How to evaluate LLMs before production

These are the lessons we learned evaluating LLMs for real-world secret scanning.




Visit: [How to evaluate LLMs before production - The GitHub Blog](#)

3

Read It · Evaluation

Benchmark Retrieval Patterns

For Copilot Studio and Foundry IQ



Compare Copilot Studio and Foundry IQ retrieval patterns with a reproducible evidence framework that keeps latency boundaries, quality gates, and cost models separate—so you can make architecture decisions that are traceable, challengeable, and repeatable.

MICROSOFT FOUNDRY · BENCHMARKING · PART 2

Benchmarking

Retrieval patterns for Copilot Studio & Foundry IQ

An evidence system, not a leaderboard —
measure the boundary · separate the cost lanes · gate on deterministic checks




The scatter plot shows a positive correlation between p95 latency and quality. A dashed box highlights a region of high quality and low latency, with a label '1 quality' and an arrow pointing to the top-left corner of the box. The x-axis is labeled 'p95'.

- 1 Boundary first**
Front door vs deployed
Time what you ship
Never subtract boundaries
- 2 Two cost lanes**
Credits vs per-token
Priced on their own meter
Never summed
- 3 Gates first**
Deterministic, then judges
Reproducible & auditable
Judges are supplemental
- ✓ One workload**
A · A2 · B · C · Hosted
Compared with evidence
Not a single number

MICROSOFT FOUNDRY BLOG 15 MIN READ

Azure Cosmos DB

Data Tools for AI Agents



Azure Cosmos DB extension for Visual Studio Code provides a way to see the editor context, understand the actual schema, and run a query with the developer's permission.

Azure Cosmos DB in the Agentic Era: Data Tools for Developers and AI Agents



Sajeetharan Sinnathurai
Principal Program Manager

[Anthropic's 2026 Agentic Coding Trends Report captures](#) an important tension in how developers use AI. Engineers report using AI in roughly 60% of their work, yet say they can fully delegate only 0–20% of their tasks. AI is becoming a constant collaborator, but it still needs setup, supervision, validation, and human judgment.

5

Try It · Azure Cosmos DB

Visit: [Azure Cosmos DB in the Agentic Era: Data Tools for Developers and AI Agents - Azure Cosmos DB Blog](#)



Reactor

MODEL MONDAYS S4:E05

Claude Models in Microsoft Foundry Tool Integrations in Practice



Hosted By: Gustavo Cordido

SEP 14, 2026 · 1:30-2:30 PM ET

Amanda
Wong

Applied AI
Anthropic



Haoran
Cheng

Product
Manager,
Microsoft



WATCH THE
LIVESTREAM



Mon 1:30pm ET
Register Here

<https://aka.ms/model-mondays/RSVP>


JOIN THE
COMMUNITY



Fri 1:30pm ET
Join The Discord

<https://aka.ms/model-mondays/Discord>

STAR US ON
GITHUB



Explore Blogs,
Recaps & Schedules

<https://aka.ms/model-mondays>