



Reactor

MODEL MONDAYS S4:E05

Claude Models in Microsoft Foundry Tool Integrations in Practice



Hosted By: Gustavo Cordido

SEP 14, 2026 · 1:30-2:30 PM ET

Amanda
Wong

Applied AI
Anthropic



Haoran
Cheng

Product
Manager,
Microsoft





Reactor

LAST WEEK IN FOUNDRY

News Highlights

Sharmila Chockalingam

Director, AI Product Marketing Microsoft




1

Try It · GPT models

New GPT Models

GPT-Live-1 and GPT-Image 2.5



Create multimodal applications with OpenAI models



**From conversations to creation:
building multimodal AI applications with
GPT-Live-1 and GPT-Image 2.5 models**

MICROSOFT FOUNDRY BLOG 8 MIN READ

Create multimodal applications with OpenAI models in Microsoft Foundry


 **Dave Jacobs**  MICROSOFT
Sep 10, 2026

Introducing GPT-Live-1 and GPT-image-2.5 models


Visit: [Create multimodal applications with OpenAI models in Microsoft Foundry | Microsoft Community Hub](#)

Azure Speech Best Practices

Call center use cases



Sample rate, chunking, speaker labels, and redaction usually are the details that separates two great speech products



MICROSOFT FOUNDRY BLOG 6 MIN READ

Voice to Chat Best Practice for Call Centre Use Cases



NaufalPrawironegoro MICROSOFT

Sep 10, 2026

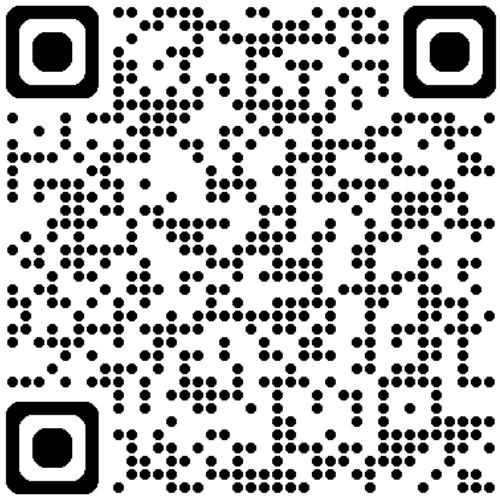
Speech projects rarely fail because someone picked the wrong model. They fail on the twenty minutes of audio nobody listened to before launch. The Open ASR Leaderboard now compares 86 systems across 12 datasets and standardises both word error rate and inverse real time factor, and even the top entry sits at 5.42 percent average word error rate while OpenAI Whisper large v3 sits at 7.44

2

Try It · Azure Speech

Azure AI Speech

Better Multilingual
Accuracy



Azure AI Speech helps developers build more accurate and capable voice experiences

Announcing Azure AI Speech LLM 2607: Better Multilingual Accuracy, Easier Customization



Rena Liu
Product Manager II

Azure AI Speech continues to deliver regular improvements to LLM Speech throughout the year, helping developers build more accurate and capable voice experiences. We're announcing **LLM Speech 2607**, our latest model update, which brings improvements in multilingual recognition and mixed-language audio inputs, along with an enhanced experience for the phrase list feature.

Table of contents

Key Advancements in Model Quality and Accuracy

New Enhancement: Phrase Lists

Example request including phrase list

Available through Fast API and Real-Time API


Get started

3

Try It · Azure Speech

Azure for Government

Microsoft Foundry
Agent Service



Accelerate mission outcomes, empower their workforce
and deliver services more effectively

Advancing trusted AI innovation with Agents and new Azure OpenAI models in Azure Government



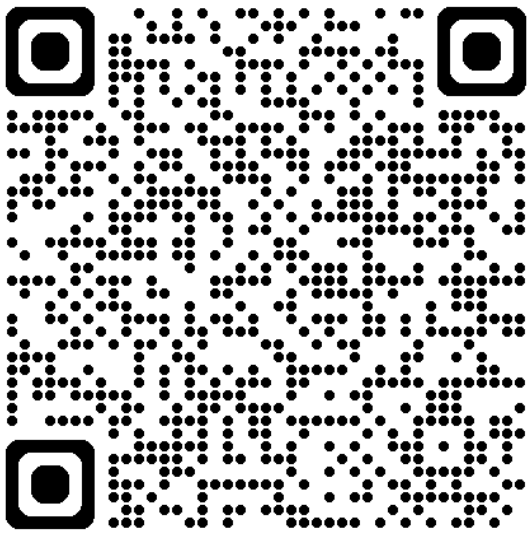
Douglas Phillips
President and Chief Technology Officer

Today, we are announcing expanded AI capabilities available across the Microsoft Azure for US Government portfolio. The Microsoft Foundry Agent Service is now available in Azure Government, giving organizations managed infrastructure on which to build, deploy, and operate AI systems that can reason over authorized information, use approved tools, and help complete multi-step work. Additionally, GPT-5.6 Sol, GPT-5.6 Terra, and GPT-5.6 Luna are now available through Microsoft.

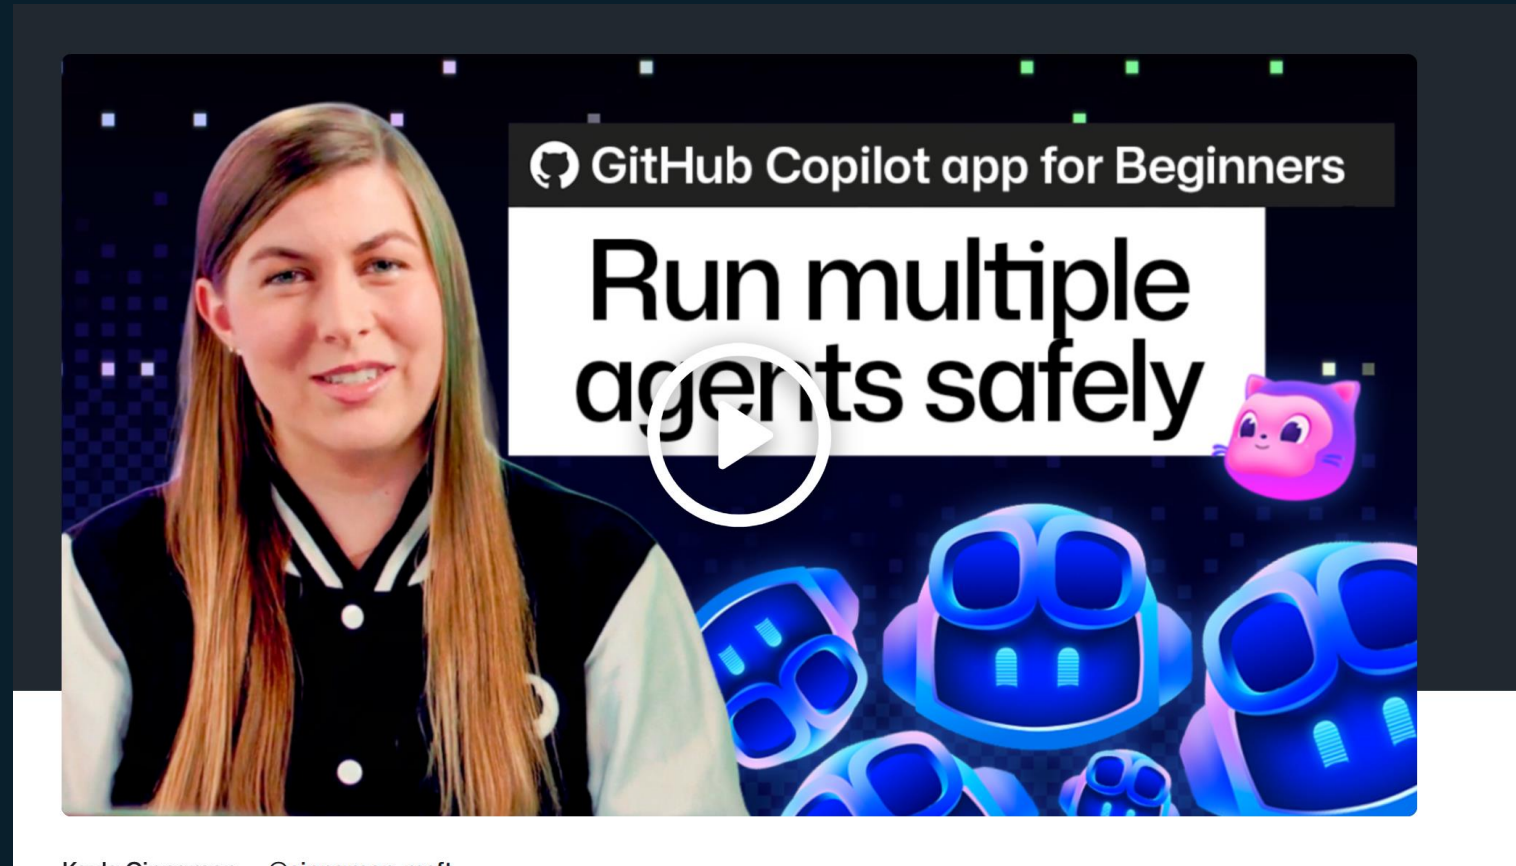
1 Try It · Microsoft Foundry
Agent Service

Github Copilot App

Run multiple AI Agents
at once



In GH Copilot App agents work separately and don't interfere with each other allowing you to get more done in less time.



5

Try It · GH Copilot App



Reactor

Claude Models on Foundry

S4:E05 – SEP 14, 2026

Amar Badal

Senior Manager,
Microsoft Foundry AI Models Ecosystem



THE ENTERPRISE PLATFORM FOR AGENTIC AI



Microsoft Foundry

One Complete Platform That Adapts To Your Evolving Business Needs

Build in GitHub (Code & Copilot) → **Run and optimize** in Foundry (Deploy & Scale) → **Reach Users** in Teams & M365



Models



Agent Service



IQ



Tools



Machine Learning



Control Plane

Cloud



Edge

Security, compliance, and governance

Every model.
Every use case.
One platform.


Avoid model lock-in:

- Flexible model selection
- Curated for organizations
- Explore, compare, evaluate and swap models with ease

Microsoft Foundry Models



11,000+ Frontier, specialized and open models
in Microsoft Foundry Models



Build with Claude in Microsoft Foundry

1

Frontier intelligence
built to perform

Frontier models with
breakthrough
reasoning, coding,
and analysis
capabilities

2

Built-in safety and
governance

Enterprise-grade
content safety, data
privacy, and
compliance controls
built in

3

Integrated with your
Microsoft stack

Native integration
with Azure, M365,
GitHub Copilot, and
the full Microsoft
ecosystem

4

From pilot to
production, faster

Streamlined
deployment with
observability,
evaluation, and
fine-tuning to
accelerate value

Claude is GA in Microsoft Foundry



Hosted on Azure

NEW AT GA

Inference in Azure (US data zone available), operated by Anthropic.

MODELS

Haiku 4.5, Sonnet 5, Opus 4.8, Opus 5

Hosted on Anthropic

Full Claude Platform feature set, reached through the customer's Azure account.

MODELS

Full lineup, including Fable 5

BOTH OPTIONS

Deploy from the Foundry model catalog like any other model.

One family. Three models. All on Foundry.

Same API, Same Safety, Route to the right intelligence for the task. Keep margins healthy without giving up the frontier.

Opus 5

Frontier intelligence

Best for

Complex reasoning.
Long-horizon agentic
work.

\$5 / \$25 per M tokens

Sonnet 5

Balanced workhorse

Best for

Most production agentic
workloads.

\$2 / \$10 per M tokens

Haiku 4.5

Fast and cost-efficient

Best for

High-volume, latency-
sensitive work.

\$1 / \$5 per M tokens

Microsoft Foundry

haorancheng74-7827

Search with AI (36K)

New Foundry

Home

Discover

Build

Operate

Manage

Docs

Welcome, Haoran Cheng

Use a model

Manage your existing endpoints or deploy a new model.

View deployments

Build an agent

Create an AI agent by configuring instructions, tools, and knowledge without writing code.

Start building

Code an agent

Define and deploy an agent in code using Microsoft Agent Framework for full control and customization.

Start coding

API key

Project endpoint

Azure OpenAI endpoint

Model selection

Test in playground

Run prompts and see responses.

Explore models

Browse available models.

Compare models

Review benchmarks and capabilities.

claude-opus-5

Anthropic

gpt-5.6-sol

OpenAI

MAI-Image-2.5

Microsoft

DeepSeek-V4-Pro

DeepSeek

Recent work

AllAgentsModels

Name	Version	Type	Last viewed	Endpoint
claude-opus-5-2	claude-opus-5	Model	9/1/26, 9:09:33 AM	View endpoint
claude-opus-5	claude-opus-5	Model	9/1/26, 8:58:42 AM	View endpoint
Azure-AI-Content-Understanding	Azure-AI-Content-Understanding	Model	8/31/26, 8:59:55 PM	View endpoint
Azure-Content-Understanding-Read	Azure-Content-Understanding-Read	Model	8/31/26, 8:59:54 PM	View endpoint
claude-hellor-5.6	claude-hellor-5.6	Model	8/31/26, 8:59:54 PM	View endpoint

TermsPrivacy

A portrait of Haoran Cheng, a man with dark, wavy hair, wearing a dark blue t-shirt. He is looking directly at the camera with a neutral expression. The background is a blurred indoor setting with light-colored walls and a window.

Haoran Cheng
Product Manager at Microsoft

Announcing new Claude capabilities in Foundry

All on the Foundry endpoint



Tool search

Picks the right tool from your library



Web search

Finds current sources



Web fetch

Pulls the full document



MCP connector

Calls the customer's own systems, CRM, ticketing, internal APIs




Structured outputs

Returns the result in the exact schema the system expects

Your systems

The outside world

The output



DEMO

Field Technician Agent

DEMO

Field Technician Agent

ANTHROPIC



Amanda Wong
Applied AI at Anthropic

Governance

What the agent is allowed to do

8 tools | 6 enabled

```
"configs": {  
  "delete_asset": {"enabled": False},  
  "update_billing": {"enabled": False}  
}
```

Two config lines. Everything else stays on.

Scoping is the first control. The MCP server still enforces auth.

What leaves Azure

AZURE

Claude in Foundry + your MCP server
Prompts and completions stay here

PUBLIC WEB

Fabrikam Manufacturer's notice
Explicit egress, domain-restricted

ANTHROPIC

Usage metadata only
No prompts, no completions



Long-running agents

Tools for what the agent reaches.
Roles for how it decides.

What it reaches



MCP
connector



Web
fetch



Structured
outputs



Web
search



Tool
search

How it decides



Planner
Decides what
to do next



Generator
Does the
work



Evaluator
Judges whether
it is done

Same Messages API. Same Foundry endpoint.
Three system prompts.

From AI-assisted to AI-driven coding

Two problems stand between an autocomplete and an agent you can hand real work to.



Knowing when it's done

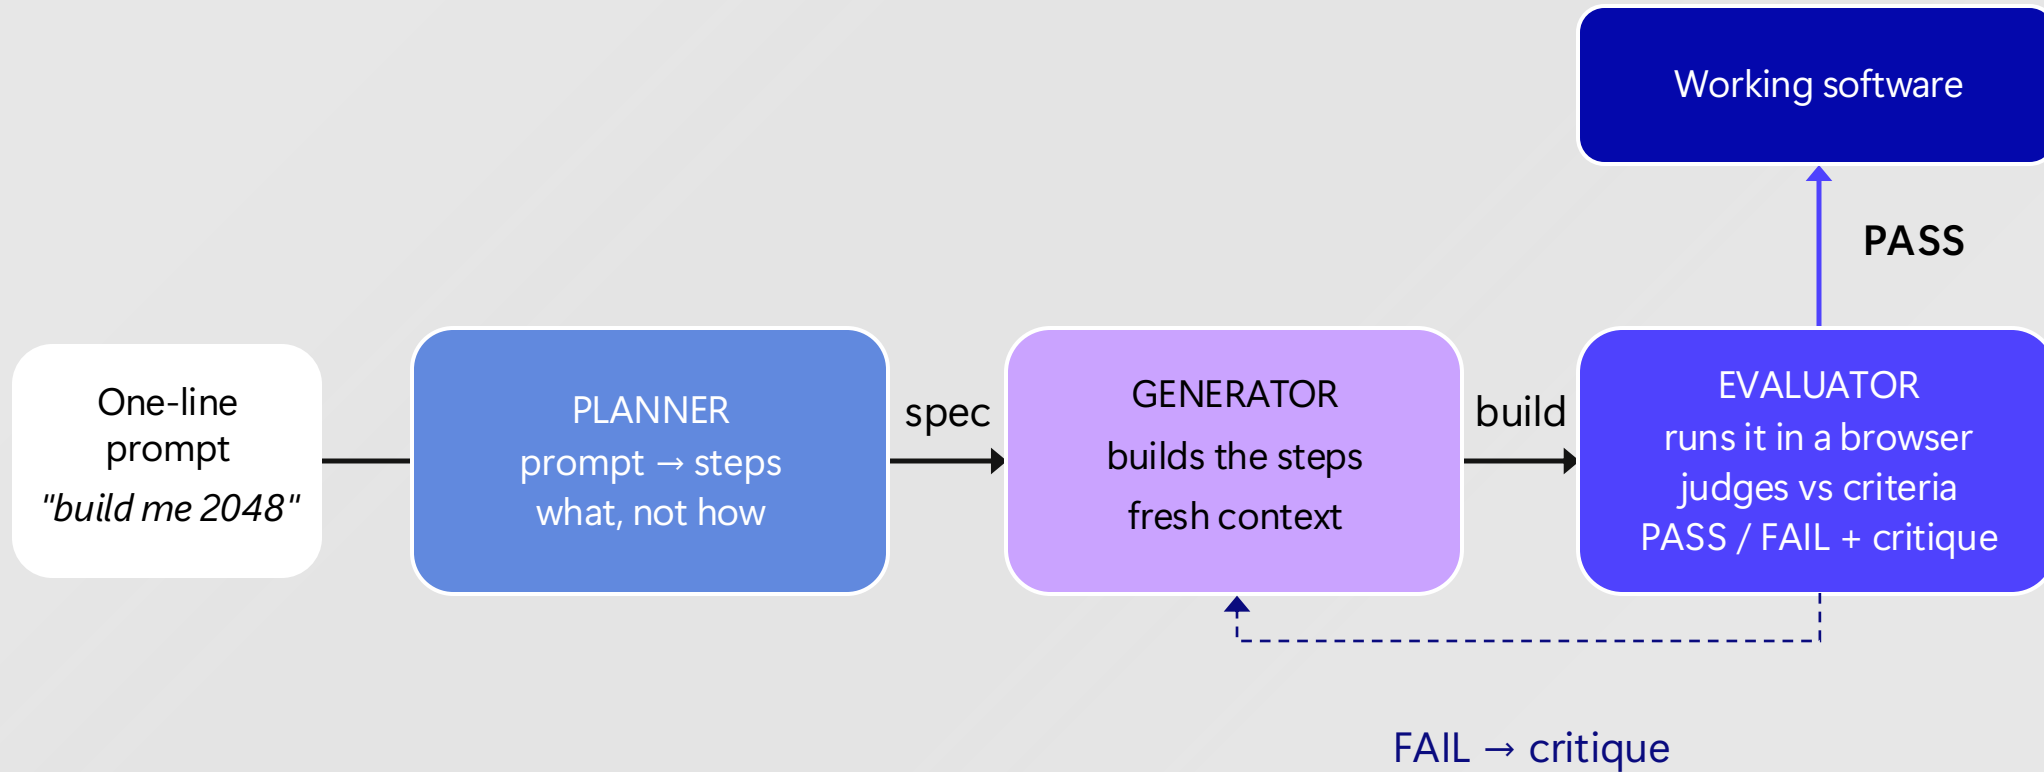
Without an explicit definition of done, agents wrap up at the first thing that compiles, not the thing the user asked for.



Knowing whether it's done well

An agent that grades its own homework will pass itself. It accepts results it shouldn't, and you only find out in production.

Three agents, one loop



All three roles are Claude Opus 5 on Microsoft Foundry – one endpoint, three system prompts.

Frontier intelligence, on your existing Azure commitment

Token spend is a design choice — model tier, caching, and batching move it by an order of magnitude.

CCU

One Azure invoice. Claude Consumption Units bill through Azure and decrement your **MACC** — no separate vendor approval.¹

90%

Cheaper cached input. Cache reads bill at 0.1× the input rate — prompts, docs, and codebases are paid for once.²

50%

Off with batch. Every model, half price for work that doesn't need a real-time response.²

RIGHT-SIZE THE MODEL TO THE TASK

MODEL	INPUT / 1M	OUTPUT / 1M
Claude Haiku 4.5 Classification, extraction, routing, high volume	\$1.00	\$5.00
Claude Sonnet 5 The production default — most agentic workloads	\$2.00	\$10.00
Claude Opus 5 Hard reasoning, long-horizon agents, research	\$5.00	\$25.00

Anthropic list token rates; Foundry converts token usage to CCU at a fixed published CCU price.
Confirm current rates on the Claude offer in Azure Marketplace before quoting.¹

THE THREE LEVERS

- 1 Tier by complexity.** Most production traffic is not frontier-hard. Routing the easy 70% to Haiku cuts spend first.
- 2 Cache the constant.** Reused context — prompts, policy docs, schemas — is written once and read at 0.1×.
- 3 Batch the async.** Evals, backfills, and bulk document processing carry no latency requirement.

Where to go next

Get started



Deploy Claude

From the Foundry model catalog, like any other model.



GA announcement

What's new at GA: hosting options, models, billing.

Build



The five tools, with code

Working samples of capabilities discussed today.



Tool docs

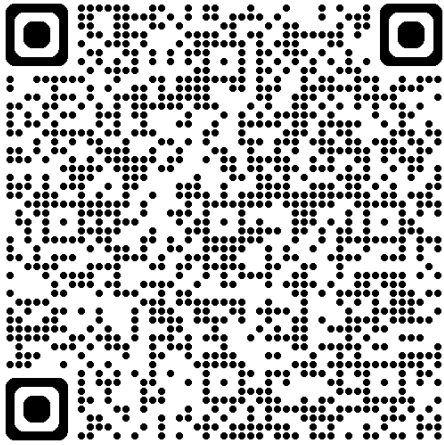
Parameters, limits, and schemas for all five.

Go deeper




Long-running agent harness

Planner, generator, evaluator, in depth.



Claude Fable 5.1
now available in
Microsoft Foundry





Claude Opus 5 is available
in Microsoft Foundry

Join Microsoft and Anthropic in San Francisco for a half-day workshop bringing technical minds together to get hands-on with Claude models in Microsoft Foundry.

This workshop is designed to accelerate your ability to ideate, optimize, and deploy AI solutions in real-world scenarios.

Request an
invite to join.

Oct 8, 2026
9:30am –
3:30pm PT




Event: <https://luma.com/p6m4n0g7>

More Workshops: <https://aka.ms/model-mastery>

2

Model Mastery Workshop



Foundry Fridays AMA

AGENT INSIGHTS & AGENT OPTIMIZER



Friday Sep 25, 1:30-2:00 pm ET

<https://aka.ms/model-mondays/chat>



Charles Kim

Product Lead / Core AI @Microsoft



Reactor

MODEL MONDAYS S4:E06

Spotlight on Azure OpenAI



Hosted By: Nitya Narasimhan

Sep 21, 2026 · 1:30-2:30 PM ET



WATCH THE
LIVESTREAM



Mon 1:30pm ET
Register Here

<https://aka.ms/model-mondays/RSVP>

JOIN THE
COMMUNITY



Fri 1:30pm ET
Join The Discord

<https://aka.ms/model-mondays/Discord>

STAR US ON
GITHUB



Explore Blogs,
Recaps & Schedules

<https://aka.ms/model-mondays>