



Reactor

MODEL MONDAYS S4:E06

Spotlight on Azure OpenAI



Hosted By: Nitya Narasimhan

Sep 21, 2026 · 1:30-2:30 PM ET



Ananya Bishnoi

AI Product Manager Lead / Microsoft



Reactor

LAST WEEK IN FOUNDRY

News Highlights


Sharmila Chockalingam

Director, AI Product Marketing Microsoft



Gartner MQ

Microsoft named a
leader for Distributed
Hybrid Infrastructure



With Azure Local and Azure Arc, organizations can manage infrastructure across datacenters, edge locations, multi-cloud environments, and sovereign deployments through a consistent Azure foundation.

September 16 • 3 min read

Microsoft recognized as a Leader in the 2026 Gartner® Magic Quadrant™ for Distributed Hybrid Infrastructure

By [Meena Gowdar](#), Partner, Product Management

Listen to this post




Gartner® names
Microsoft a Leader in

Distributed Hybrid
Infrastructure

Visit: [Microsoft recognized as a Leader in the 2026 Gartner® Magic Quadrant™ for Distributed Hybrid Infrastructure | Microsoft Azure Blog](#)

Microsoft Foundry Agents

Agent to agent collaboration has become easier in Foundry



Build multi-agent systems where agents can specialize, publish their skills, and securely collaborate across service boundaries.



BLOG POST

A2A Endpoints and A2A Tool in Microsoft Foundry agents

Build more capable agents by connecting with other agents using the Agent-to-Agent (A2A) protocol.



Collaborate

Agents discover and communicate securely



Extend

Use A2A Tool to leverage remote agent capabilities




Scale

Create modular, interoperable agent ecosystems



A2A Endpoints

Expose your agent as an A2A endpoint so other agents can discover and interact with it.



A2A Endpoint

- Discoverable
- Secure
- Standards-based (A2A Protocol)



A2A Tool

Use the A2A Tool to call skills and capabilities exposed by remote agents.



A2A Tool



- Plug & Play
- Reuse capabilities
- Faster innovation

Response


Visit: [A2A Endpoints and A2A Tool in Microsoft Foundry agents | Microsoft Community Hub](#)

2

Try It · Foundry Agents

Local and contained AI

Foundry Local, Azure
ML, Azure Local




Which execution location best satisfies the business, data, security, performance, and operational requirements of each workload?




Microsoft Foundry Agents

Integration patterns for Microsoft Foundry Agents to Enterprise data platforms



How the agent should access that data while preserving the right security, governance, semantics, freshness and system-of-record boundaries.




1


Try It · Foundry Agents

Github Podcast

Hot takes from the AI
world



A good hot take gives you something sharp enough to question. The questions are where you find the useful ideas.



5 Listen · Github Podcast

Visit: [Should you read the code, is RAG dead, and did Skills kill MCP? - The GitHub Blog](#)



Reactor

Model Mondays

<https://aka.ms/model-mondays>

JOIN THE COMMUNITY

<https://aka.ms/model-mondays/discord>





Reactor

WHAT'S NEW WITH GPT?

Spotlight on Azure OpenAI

Ananya Bishnoi

Technical Program Manager, Microsoft
Foundry




What's new in Microsoft Foundry

GPT-5.6 Sol, Terra, Luna • GPT-6 Astra

A guide to capabilities, trade-offs, and model selection

Ananya Bishnoi | September 2026



A broader frontier portfolio

Four models span one decision space: maximum intelligence, balanced economics, speed, and throughput.

GPT-6 Astra	Frontier/Hardest problems	Best overall capability
GPT-5.6 Sol	Deep reasoning	Frontier for 5.6
GPT-5.6 Terra	Balanced enterprise	Performance-cost sweet spot
GPT-5.6 Luna	Fast and efficient	Low-latency, high-volume, cost efficient

Shared foundation

Text + image input

Text output

Up to ~1M context, 1.05M for Astra

Tool-enabled workflows

GPT-6 Astra

Frontier agentic intelligence for open-ended knowledge work and execution

Built for tasks where AI needs to go beyond answering questions and actually reason, plan and execute multi-step work.

INTELLIGENCE + AGENCY

● Advanced reasoning

Performs deep, multi-step reasoning for complex problem solving and decision support.

● Agentic workflows

Plans, executes, and coordinates multi-step tasks using tools and long-running workflows.

● Long-context understanding

Supports up to 1.05M tokens for large documents, repositories, and knowledge bases.

● Advanced coding and engineering

Excels at code generation, debugging, software design, and technical problem solving.

● Connected intelligence

Enhanced browsing, retrieval, connectors, and file search for discovery and synthesis.

● Computer use and automation

Supports computer-use scenarios and tool-enabled workflow automation.

Selection note: Astra is the strongest fit when the model must plan, use tools, adapt, and deliver a polished work product.

GPT-5.6 Sol

Flagship reasoning for the hardest, most compute-intensive workloads

DEPTH + ORCHESTRATION

● Advanced reasoning

Performs deep, multi-step reasoning for complex problem solving and decision support.

● Multi-agent orchestration

Coordinates specialized subagents that collaborate on sophisticated tasks.

● Long-context understanding

Supports up to 1M tokens for large documents, code repositories, and knowledge bases.

● Advanced coding and engineering

Excels at code generation, debugging, software design, and technical analysis.

● Research and scientific workflows

Strong performance on complex research, biology, and scientific reasoning tasks.

● Cybersecurity analysis

Designed for vulnerability research and long-horizon security investigations.

Selection note: Use Sol when the problem is genuinely hard enough to justify premium reasoning and compute.

GPT-5.6 Terra

Balanced reasoning, tool use, and economics for everyday enterprise AI

BALANCE + VALUE

● Cost-efficient intelligence

Competitive performance at lower cost for scaling AI across the enterprise.

● Balanced reasoning

Flexible trade-off between speed, cost, and depth as reasoning effort increases.

● Tool use and execution

Uses tools and external systems to complete real-world business workflows.

● Long-context understanding

Supports up to 1M tokens for documents, knowledge bases, and multi-step workflows.

● Enterprise scale

Optimized for reliable, high-volume deployment across broad business scenarios.

● Robust safety protections

Safeguards are calibrated to model capabilities to help detect and mitigate risk.

Selection note: Terra is the first model to test for broad production workloads when both quality and cost matter.

GPT-5.6 Luna

Fast, cost-efficient inference for real-time and high-volume experiences

SPEED + SCALE

● Low latency

Fast response times for real-time applications and interactive experiences.

● Cost-efficient at scale

Strong capability at the lowest cost in the GPT-5.6 series for high-volume workloads.

● Scalable reasoning

Performance improves as reasoning effort increases on more complex tasks.

● Long-context understanding

Supports up to 1M tokens for large documents, conversations, and knowledge bases.

● Multilingual workloads

Optimized for high-volume language experiences across markets and use cases.

● Built-in safeguards

Safety mitigations are calibrated to Luna's capabilities and deployment profile.

Selection note: Start with Luna when latency and unit economics dominate, then move up only if quality testing shows a gap.

Customer model selection

Start with the workload constraint, then validate quality, latency, and cost on representative data.

CUSTOMER QUESTION

START WITH

WHY

Need the strongest end-to-end agentic work?

GPT-6 Astra

Plan, connect, automate, and produce polished outputs

Need maximum 5.6 reasoning depth?

GPT-5.6 Sol

Complex analysis, coding, research, orchestration

Need a strong enterprise default?

GPT-5.6 Terra

Balanced quality, tools, scale, and cost

Need low latency or very high volume?

GPT-5.6 Luna

Real-time and cost-sensitive applications

Pricing comparison

Paygo list pricing in USD per 1 million tokens; long context begins above 272K input tokens


Model	Context	Input	Output	Customer fit
GPT-6 Astra	Short	\$10.00	\$50.00	Frontier agentic work
	Long	\$20.00	\$75.00	
GPT-5.6 Sol*	Short	\$4.00*	\$20.00*	Deep reasoning
	Long	\$10.00	\$45.00	
GPT-5.6 Terra	Short	\$2.00	\$12.00	Balanced enterprise default
	Long	\$4.00	\$18.00	
GPT-5.6 Luna	Short	\$0.20	\$1.20	Low latency and high volume
	Long	\$0.40	\$1.80	

Compare cost per completed task using representative workloads, not token price alone.

Thank you!

Explore the Azure OpenAI Collection on Microsoft Foundry

gpt-6-astra
gpt-5.6- family
gpt-live-1
gpt-image-2.5 family



Microsoft Foundry

Enterprise Discover API reference Pricing Sign in

Catalog / Models

Models Agents Tools Solution templates

Frontier, open-source, and partner models in one place.

Search Publisher (1) Inference tasks Supported features Region All filters

70 models

Publisher: OpenAI Clear filters

 gpt-6-astra chat-completion	 gpt-5.6-sol chat-completion	 gpt-5.6-luna chat-completion
Input type Output type Context window Token limits	Input type Output type Context window Token limits	Input type Output type Context window Token limits
text, image text 1050k 128k output	text, image text 1050k 128k output	text, image text 1050k 128k output
Generally available (GA)	Generally available (GA)	Generally available (GA)


 gpt-5.6-terra chat-completion	 gpt-realtime-2.1 audio-generation	 gpt-realtime-2.1-mini audio-generation
Input type Output type Context window Token limits	Input type Output type Context window Token limits	Input type Output type Context window Token limits
text, image text 1050k 128k output	audio, text, image audio, text 128k 4.096k output	audio, text, image audio, text 128k 4.096k output
Generally available (GA)	Preview	Preview

View: <https://ai.azure.com/catalog/models?publisher=openai>

Turn open-ended goals
into action

Computer use across
applications

Enterprise controls for
agentic work



GPT-6-Astra – Frontier Intelligence for work, now generally available in Microsoft Foundry

Computer use across applications


Astra's computer-use capabilities are designed to work across familiar applications, including workflows without dedicated APIs. It can interpret on-screen information and interact with approved interfaces to support tasks such as updating records, navigating development tools, testing software, and assembling results into reports. OpenAI reports state-of-the-art results on selected computer-use evaluations; performance varies by task, tools, configuration, and safeguards.

Capability this direct demands containment. Content displayed in an application may be incomplete, misleading, or designed to influence an agent's behavior. Foundry helps customers define access, approvals, and monitoring, and design workflows with scoped credentials, approved resources, human checkpoints for consequential actions, and activity records aligned to their risk requirements.

Enterprise scenarios we're seeing

- **Software engineering:** Astra can reproduce complex bugs, investigate likely causes, propose fixes, and prepare changes for developer testing and review.
- **Business intelligence:** Astra can build and refine dashboards in Power BI, helping analysts compare data, identify trade-offs, and prepare insights to share.
- **Professional work:** Astra can produce documents, spreadsheets, and presentations that follow existing templates and business standards, creating polished artifacts for expert review.
- **Application workflows:** Astra can support tasks such as updating customer records, processing forms, testing websites, and working through approved interfaces where dedicated APIs are limited.

the GPT-5.6 series gives organizations the flexibility to match model capability, cost, and performance requirements to specific business scenarios, rather than forcing every workload onto a single model.



An agent is only as capable as the model reasoning behind it. The GPT-5.6 series is generally available in Microsoft Foundry Models and Microsoft Foundry Agent Service. Use the latest frontier models and right-size the model to the agentic task.



View: <https://azure.microsoft.com/en-us/blog/gpt-5-6-now-available-in-microsoft-foundry/>

From conversations to creation

What are GPT-Image-2.5 Flare and GPT-Image-2.5 Sunburst?


The GPT-Image-2.5 family gives developers two options for building image generation and editing into user-facing applications. **GPT-Image-2.5 Flare** is a fast, high-throughput model designed for most production workloads. As the smaller model, Flare delivers improved image quality, editing accuracy, and responsiveness for content creation, marketing assets, product experiences, visual search, and image generation at scale. **GPT-Image-2.5 Sunburst** is a premium model optimized for creative workflows that require greater precision and control. Sunburst is designed for high-fidelity campaign assets, product imagery, and other polished visual experiences where quality takes priority over generation speed. Together they bring:

- **Significantly faster generation and editing.** GPT-Image-2.5 Flare delivers higher-quality images than GPT-Image-2 at 50% lower latency, helping keep users in a live, conversational iteration loop.
- **More accurate editing.** Update targeted elements while maintaining the rest of the image.
- **Stronger multi-turn editing.** Refine images through successive instructions without losing consistency or image quality.
- **Better instruction following.** Interpret complex visual instructions, layouts, and stylistic directions more accurately.
- **Two models for different workflows.** Use Sunburst for premium creative and editing workflows and Flare, the smaller model, for faster iteration.

See it in action:


Below is a series of images that demonstrates the improvements in multi-turn editing with GPT-images 2.5 Sunburst. The sequence shows the model responding to a series of editing instructions over multiple turns. As changes accumulate, the model maintains visual consistency, preserves details from earlier edits, and incorporates new requests without degrading image quality. These examples highlight how GPT-images 2.5 Sunburst enables more reliable iterative editing workflows, making it easier to refine assets through a sequence of targeted updates.

Prompt: A pair of running shoes placed on a forest trail at sunrise, sitting on packed dirt scattered with fallen leaves and pine needles. Soft golden sunlight filters through tall trees in the background, casting long shadows across the path. Morning mist hovers low between the trunks. Shallow depth of field, the shoes in sharp focus, photorealistic, natural color palette, cinematic lighting.



For developers and creators, this means less time recreating assets and more time refining them, with the confidence that each edit will build on the last.

Building multimodal AI applications with GPT-Live-1 and GPT-Image 2.5 models on Microsoft Foundry





Reactor

MICROSOFT FOUNDRY DEMOS

Hands-on With Azure OpenAI


Nitya Narasimhan, PhD

Senior AI Advocate, Microsoft



Want to build your intuition for what a new model release does?

Explore the repo for sample notebooks and scenarios you can extend



Find the right model for your next agentic task.

Microsoft Foundry [catalog](#) Model Mondays [livestream](#) Model Mastery [workshops](#) License [MIT](#)

[Changelog](#) · [Capsules](#) · [Quickstart](#) · [Reference](#)

The [Microsoft Foundry catalog](#) has thousands of models from Anthropic, Microsoft, OpenAI, xAI, Hugging Face, Meta, Mistral, Cohere, and NVIDIA - and new releases land almost daily. Keeping up is one thing. Knowing what a release actually does *differently*, and whether it's the right pick for your agent, is another.

That's what this repo is for:

1. **Catching up?** The [CHANGELOG](#) has every announcement in one place.
2. **Want to try one?** Browse [models/](#) for a release *capsule* - a runnable notebook that shows you what it does.
3. **Can't find the one you need?** [Open an issue](#) and we'll backfill it.

Code-first samples for new releases: <https://aka.ms/model-releases>



Foundry Fridays AMA

AGENT INSIGHTS & AGENT OPTIMIZER



Friday Sep 25, 1:30-2:00 pm ET

<https://aka.ms/model-mondays/chat>



Charles Kim

Product Lead / Core AI @Microsoft



Reactor

MODEL MONDAYS S4:E06

Mistral OCR-4 & Building AI Agents in Microsoft Foundry



Hosted By: Amy Boyd

SEP 28, 2026 · 1:30-2:30 PM ET



**Victor
Rojo**

Technical Partner
Enablement Lead,
Mistral AI



**Peyman
Mohajerian**

Partner Solution
Architect, Mistral AI

WATCH THE
LIVESTREAM



Mon 1:30pm ET
Register Here

<https://aka.ms/model-mondays/RSVP>


JOIN THE
COMMUNITY



Fri 1:30pm ET
Join The Discord

<https://aka.ms/model-mondays/Discord>

STAR US ON
GITHUB



Explore Blogs,
Recaps & Schedules

<https://aka.ms/model-mondays>